

World Life Expectancy Project (Data Cleaning)

Look at the data after having imported them into MySQL Workbench.

```
SELECT *  
FROM world_life_expectancy  
;
```

Even though Alex separates data cleaning from exploratory data analysis,
he says that in the real world both are likely to be performed at the same time.

Each row has a row ID, so duplicates can be determined by combining values from two different columns,
kinda like composite primary keys in the context of database administration. In the case of this project,
the combination of country and year can help identify unique rows since there should be one row
(i.e., a combination of values) for each combination of country and year.

```
SELECT Country, Year, CONCAT(Country, Year)  
FROM world_life_expectancy  
;
```

Everything in the column CONCAT(Country, Year) should be unique; however,
Alex quickly found a duplicate (two instances of Zimbabwe2019).

Run the same query as before but with a count of instances of CONCAT(Country, Year).
Reminder: Aggregate functions in the SELECT statement require a GROUP BY statement
later on in the same query.

```
SELECT Country, Year, CONCAT(Country, Year), COUNT(CONCAT(Country, Year))  
FROM world_life_expectancy  
GROUP BY Country, Year, CONCAT(Country, Year)  
;
```

Run the same query as before but with a HAVING statement this time
to identify counts under COUNT(CONCAT(Country, Year)) greater than 1.
Reminder: Aggregate functions in the SELECT statement require a GROUP BY statement
later on in the same query.

```
SELECT Country, Year, CONCAT(Country, Year), COUNT(CONCAT(Country, Year))  
FROM world_life_expectancy  
GROUP BY Country, Year, CONCAT(Country, Year)  
HAVING 1 < COUNT(CONCAT(Country, Year))  
;
```

Look at the data once again and take note of the column Row_ID.

```
SELECT *  
FROM world_life_expectancy  
;
```

Generate row numbers over partitions.

```
SELECT Row_ID,
```

```

        CONCAT(Country, Year),
        ROW_NUMBER() OVER(PARTITION BY CONCAT(Country, Year) ORDER BY CONCAT(Country, Year)) AS
Row_Num
FROM world_life_expectancy
;

```

In order to filter on the result of the above query,
the above query needs to be used as a subquery.

```

SELECT *
FROM (
        SELECT Row_ID,
        CONCAT(Country, Year),
        ROW_NUMBER() OVER(PARTITION BY CONCAT(Country, Year) ORDER BY CONCAT(Country, Year)) AS
Row_Num
        FROM world_life_expectancy
) AS row_table
WHERE 1 < Row_Num
;

```

Row IDs that appear in the above query will become the basis
for deleting duplicates in the world life expectancy table.

Alex goes through the steps of creating a backup of the world life expectancy table
(named world_life_expectancy_backup) in order to demonstrate best practices
when doing data cleaning in the real world.

The above query will be used in the WHERE statement of the next query
in order to identify which rows to delete. However, instead of SELECT * like
in the first subquery, SELECT Row_ID will be used instead.
I'm going to first use a SELECT * instead of going straight to a DELETE statement
just to confirm that I'm about to delete what I want to delete.
(This is just a me thing. Sometimes bad luck does happen)

```

SELECT *
FROM world_life_expectancy
WHERE Row_ID IN (
        SELECT Row_ID
        FROM (
                SELECT Row_ID,
                CONCAT(Country, Year),
                ROW_NUMBER() OVER(PARTITION BY CONCAT(Country, Year) ORDER BY
CONCAT(Country, Year)) AS Row_Num
                FROM world_life_expectancy
        ) AS row_table
        WHERE 1 < Row_Num
)
;

```

The results of the above query look good, so I'll go back to following along

with what Alex is doing in the video.

```
DELETE
FROM world_life_expectancy
WHERE Row_ID IN (
    SELECT Row_ID
    FROM (
        SELECT Row_ID,
        CONCAT(Country, Year),
        ROW_NUMBER() OVER(PARTITION BY CONCAT(Country, Year) ORDER BY
CONCAT(Country, Year)) AS Row_Num
        FROM world_life_expectancy
    ) AS row_table
    WHERE 1 < Row_Num
)
;
```

Run the following query. Nothing should show up.

```
SELECT *
FROM (
    SELECT Row_ID,
    CONCAT(Country, Year),
    ROW_NUMBER() OVER(PARTITION BY CONCAT(Country, Year) ORDER BY
CONCAT(Country, Year)) AS Row_Num
    FROM world_life_expectancy
) AS row_table
WHERE 1 < Row_Num
;
```

Alex points out that figuring out how to remove duplicates in a dataset may not be intuitive.

Look at the data once again to prepare for the next phase of the data cleaning process.

```
SELECT *
FROM world_life_expectancy
;
```

Look for the rows where the status of the country is blank.

The word NULL doesn't appear in queries for that column,

so the value is likely just a blank/empty string

```
SELECT *
FROM world_life_expectancy
WHERE Status = ""
;
```

Look at the statuses among the countries that contain at least one blank

for their development status

Alex uses the <> combination for inequality; I choose the != combination.

The outcome should be the same.

```
SELECT DISTINCT Status
```

```
FROM world_life_expectancy
WHERE Status != ''
;
```

```
# Look at countries with the 'Developing' status
SELECT DISTINCT country
FROM world_life_expectancy
WHERE Status = 'Developing'
;
```

```
# Attempt to update the world life expectancy table
# so that 'Developing' countries are set
# to 'Developing' in order to fill any blanks in status column
# for developing countries
UPDATE world_life_expectancy
SET status = 'Developing'
WHERE country IN (
    SELECT DISTINCT country
    FROM world_life_expectancy
    WHERE Status = 'Developing'
)
;
```

```
# The above query failed to execute because the table being subjected
# to the UPDATE statement can't show up in the FROM statement (of the subquery
# in this case).
```

```
# Take a different approach - perform a self-join on the world life expectancy table,
# and then try to update the table from there
# Again, Alex uses <> for inequality, while I use != instead. The outcome
# should still be the same.
UPDATE world_life_expectancy AS t1
INNER JOIN world_life_expectancy AS t2
    ON t1.Country = t2.Country
SET t1.Status = 'Developing'
WHERE t1.Status = ''
    AND t2.Status != ''
    AND t2.Status = 'Developing'
;
```

```
# Thinking about it, the statement AND t2.Status != '' in the above query feels unnecessary
# because if t2.Status = 'Developing', then it is also the case that t2.Status != '', ya know?
# However, I'll still follow along with what Alex has so that I don't inadvertently
# become confused later on.
```

```
# Run a query to confirm that the change had gone through
SELECT *
FROM world_life_expectancy
```

```
WHERE Status = "  
;
```

```
# Run a query to look at the development status  
# of the United States of America  
SELECT *  
FROM world_life_expectancy  
WHERE Country = 'United States of America'  
;
```

```
# Perform a self-join on the world life expectancy table again,  
# and then try to update the table from there once more
```

```
# Prior to attempting the same update as Alex,  
# I'm going to just run an INNER JOIN  
# of the world life expectancy with itself on country  
# and then just look at the rows for the United States of America  
# in order to understand how the previous update worked.  
SELECT *  
FROM world_life_expectancy AS t1  
INNER JOIN world_life_expectancy AS t2  
ON t1.Country = t2.Country  
WHERE t1.Country = 'United States of America'  
;
```

```
# So it looks like performing such an INNER JOIN produced  
# what looks like the result of a CROSS JOIN. The output window  
# said that 256 row(s) had been returned, and the square root of 256 is 16  
# according to a Google search.
```

```
SELECT Country, COUNT(Country)  
FROM world_life_expectancy  
WHERE Country = 'United States of America'  
;
```

```
# The above query returns 16 for the number of rows for the United States of America,  
# which matches the square root of 256, so indeed it does look like performing an INNER JOIN  
# of a table with itself along a value that repeats itself across all relevant rows  
# does result in a CROSS JOIN. (Good to know, eh?)  
# Therefore, it's possible  
# for a country-year combination that's missing a development status  
# to be paired with country-year combination that isn't missing a development status.
```

```
# Again, Alex uses <> for inequality, while I use != instead. The outcome  
# should still be the same.  
# Also, Alex opts to just copy-and-paste the value 'Developed'  
# instead of typing it out like I did in order to safeguard against  
# spelling something incorrectly. That makes a lot of sense, so going forward
```

I'll try to copy-and-paste values from the table(s) when applicable.

```
UPDATE world_life_expectancy AS t1
INNER JOIN world_life_expectancy AS t2
    ON t1.Country = t2.Country
SET t1.Status = 'Developed'
WHERE t1.Status = ""
    AND t2.Status != ""
    AND t2.Status = 'Developed'
;
```

Run a query to look at the development status
of the United States of America.

```
SELECT *
FROM world_life_expectancy
WHERE Country = 'United States of America'
;
```

Look for any other rows with a blank development status.

```
SELECT *
FROM world_life_expectancy
WHERE Status = ""
;
```

Also check for any other rows with a NULL value in the Status column.
(These NULL values should've been addressed earlier,
but it doesn't hurt to check again.)

```
SELECT *
FROM world_life_expectancy
WHERE Status = NULL
;
```

Look at the table once more.

Alex says that he likes to keep copies of queries in the same order that he's thinking
through the problem, even if such copies can become redundant.

I'm inclined to do the same since doing so helps maintain the thought process,
especially when comments are added in to explicitly explain things.

```
SELECT *
FROM world_life_expectancy
;
```

A concern that I have here is that up to this point Alex didn't explicitly factor in
the possibility of a country's development status changing from 'Developing' to 'Developed'
which then introduces the possibility that a blank status could appear
between a row that has 'Developing' and a row that says 'Developed'
when ordering those rows by year.

I suppose such a concern could be something that I myself could keep track of for this guided project
or take into consideration when working to clean data in the real world.

Isolate just the rows with blank values under the life expectancy column

```
SELECT *  
FROM world_life_expectancy  
WHERE `Life expectancy` = ''  
;
```

The above query returns a row for Afghanistan and a row for Albania.

Review the full table once again.

```
SELECT *  
FROM world_life_expectancy  
;
```

Alex briefly looks over the data for Afghanistan in the full table, and
he notices that the numbers in the life expectancy column for that country are
gradually increasing, so he recommends populating the blank for that country
with the average of two numbers - the number of the year for the country before that blank
and the number of the year for the country after that blank, provided that the rows are ordered
by year for that country.
Moreover, he says that later on when he goes through exploratory data analysis (EDA), it would be a
good idea
to have each blank space populated with some a data point.
Trying to populate the blank with the average of the values above and below it when the rows are
sorted
can be tricky, but Alex says that that's just how things go for data analysts.

Focus on a few columns of the world life expectancy table.

```
SELECT Country, Year, `Life expectancy`  
FROM world_life_expectancy  
;
```

Attempt two self-joins between the world life expectancy table and itself.

Remember to specify the table for the column names in the SELECT statement.

```
SELECT t1.Country, t1.Year, t1.`Life expectancy`,  
       t2.Country, t2.Year, t2.`Life expectancy`,  
       t3.Country, t3.Year, t3.`Life expectancy`  
FROM world_life_expectancy AS t1  
INNER JOIN world_life_expectancy AS t2  
    ON t1.Country = t2.Country  
    AND t1.Year = (t2.Year - 1)  
INNER JOIN world_life_expectancy AS t3  
    ON t1.Country = t3.Country  
    AND t1.Year = (t3.Year + 1)  
;
```

Now perform the above query again, but with focus placed

on that one row with the blank value under life expectancy.

```
SELECT t1.Country, t1.Year, t1.`Life expectancy`,
```

```

        t2.Country, t2.Year, t2.`Life expectancy`,
        t3.Country, t3.Year, t3.`Life expectancy`
FROM world_life_expectancy AS t1
INNER JOIN world_life_expectancy AS t2
    ON t1.Country = t2.Country
    AND t1.Year = (t2.Year - 1)
INNER JOIN world_life_expectancy AS t3
    ON t1.Country = t3.Country
    AND t1.Year = (t3.Year + 1)
WHERE t1.`Life expectancy` = ''
;

```

Now perform the above query again, but add a column
with the average of the second and third life expectancy columns.

```

SELECT t1.Country, t1.Year, t1.`Life expectancy`,
        t2.Country, t2.Year, t2.`Life expectancy`,
        t3.Country, t3.Year, t3.`Life expectancy`,
        ROUND((t2.`Life expectancy` + t3.`Life expectancy`)/2, 1)
FROM world_life_expectancy AS t1
INNER JOIN world_life_expectancy AS t2
    ON t1.Country = t2.Country
    AND t1.Year = (t2.Year - 1)
INNER JOIN world_life_expectancy AS t3
    ON t1.Country = t3.Country
    AND t1.Year = (t3.Year + 1)
WHERE t1.`Life expectancy` = ''
;

```

Update the world life expectancy table using what had been found with the previous query
as well as what had been discovered/learned during the previous update.
Reminder: The table selected for updating can't show up in the FROM statement.
The "double self-join" will go in the UPDATE statement.
The calculated life expectancy value should only be applied to the rows
where the life expectancy is blank.

```

UPDATE (world_life_expectancy AS t1
INNER JOIN world_life_expectancy AS t2
    ON t1.Country = t2.Country
    AND t1.Year = (t2.Year - 1)
INNER JOIN world_life_expectancy AS t3
    ON t1.Country = t3.Country
    AND t1.Year = (t3.Year + 1)
)
SET t1.`Life expectancy` = ROUND((t2.`Life expectancy` + t3.`Life expectancy`)/2, 1)
WHERE t1.`Life expectancy` = ''
;

```

Look at the "double self-join" again to confirm that the update went through as intended.
SELECT t1.Country, t1.Year, t1.`Life expectancy`,


```

        t2.Country, t2.Year, t2.`Life expectancy`,
        t3.Country, t3.Year, t3.`Life expectancy`,
        ROUND((t2.`Life expectancy` + t3.`Life expectancy`)/2, 1)
FROM world_life_expectancy AS t1
INNER JOIN world_life_expectancy AS t2
    ON t1.Country = t2.Country
    AND t1.Year = (t2.Year - 1)
INNER JOIN world_life_expectancy AS t3
    ON t1.Country = t3.Country
    AND t1.Year = (t3.Year + 1)
WHERE t1.`Life expectancy` = "
;

```

No rows showed up in the result of the above query, so that's neat.

```

# Look at the world life expectancy table again
# with a focus on just a few columns
SELECT Country, Year, `Life expectancy`
FROM world_life_expectancy
;

```

```

# Check for any rows with blanks in the life expectancy column.
SELECT Country, Year, `Life expectancy`
FROM world_life_expectancy
WHERE `Life expectancy` = "
;

```

Nothing shows up - neat.

According to Alex, figuring out the calculations to replace the blank values
is not easy, but going through the steps does help with future data cleaning procedures.

I do have a concern here, though. When looking at the trend of life expectancy values
to determine how to replace a blank value, Alex only looked at the numbers for Afghanistan.
He didn't look at the values for Albania. And yet, for both Afghanistan and Albania,
he decided to calculate the average of "flanking" life expectancy values for each country to replace
their
blank values. In the grand scheme of it all, such a decision
might not negatively impact the exploratory data analysis too much later on; however, I feel like
both the values for Afghanistan and Albania should've been observed separately before deciding on
how to replace the blank values for each country. Maybe all of this is something I can keep in mind
when doing my own data cleaning in the future.

```

# Bring up the whole life expectancy table one more time.
SELECT *
FROM world_life_expectancy
;

```

It doesn't look like there's anything else to address at this point. However, Alex points out
that it's possible that more deficiencies in the data can be found when doing EDA later on,
so those can be addressed then. Also, he points out that all of these steps taken to clean this dataset
is basically how things go in the real world when it comes to cleaning any dataset.