

# World Life Expectancy Project (Exploratory Data Analysis)

# Alex says that there are two parts to exploratory data analysis, or EDA.  
# The first part is in conjunction with data cleaning and the goal is to find out  
# what's wrong with the dataset and should be fixed. The second part focuses on  
# finding insights or trends in the data that can be used in the future.

# Look at the life expectancy table and try to find anything potentially helpful.

```
SELECT *  
FROM world_life_expectancy  
;
```

# Look at how different countries did in terms of increasing (or maybe even decreasing)  
# the life expectancy of their people.

```
SELECT Country, MIN(`Life expectancy`), MAX(`Life expectancy`)  
FROM world_life_expectancy  
GROUP BY Country  
ORDER BY Country DESC  
;
```

# In the result of the above query, there were some countries  
# where the min or max life expectancy was 0, which could indicate  
# a data quality issue; therefore, the query above will be run again  
# with a filter to exclude rows that have these zeros.  
# Note: Alex uses <> for inequality. I personally use != for inequality.  
# The outcome should be the same.

```
SELECT Country, MIN(`Life expectancy`), MAX(`Life expectancy`)  
FROM world_life_expectancy  
GROUP BY Country  
HAVING 0 != MIN(`Life expectancy`)  
        AND 0 != MAX(`Life expectancy`)  
ORDER BY Country DESC  
;
```

# Run the above query again but this time include the range (i.e., max - min)  
# of the life expectancy values, and order those ranges from highest to lowest  
# Note: I used a title for the ranges that's different from what Alex uses in the video.  
# He says that these ranges are for a span of fifteen years; however, I noticed during the  
# data cleaning process that the United States has sixteen years of data in the table. Since there  
# could be some countries that have fewer years on record than other countries,  
# I'll just keep the title a little more general

```
SELECT Country, MIN(`Life expectancy`), MAX(`Life expectancy`),  
        ROUND((MAX(`Life expectancy`) - MIN(`Life expectancy`)), 1) AS  
Life_Increase_Over_All_Years_on_Record  
FROM world_life_expectancy  
GROUP BY Country  
HAVING 0 != MIN(`Life expectancy`)  
        AND 0 != MAX(`Life expectancy`)
```

```
ORDER BY Life_Increase_Over_All_Years_on_Record DESC
;
```

# Run the above query again but sort the ranges in ascending order to see which countries  
# had the smallest increase in life expectancy.

```
SELECT Country, MIN(`Life expectancy`), MAX(`Life expectancy`),  
       ROUND((MAX(`Life expectancy`) - MIN(`Life expectancy`)), 1) AS  
Life_Increase_Over_All_Years_on_Record  
FROM world_life_expectancy  
GROUP BY Country  
HAVING 0 != MIN(`Life expectancy`)  
       AND 0 != MAX(`Life expectancy`)  
ORDER BY Life_Increase_Over_All_Years_on_Record ASC  
;
```

# It looks like countries that have a high min value to begin with  
# will have small differences from the max value

# I'm noticing that the results of the EDA so far are prompted by curiosity about the data.  
# Like...first asking one question about a subset of the data, and then asking further questions  
# about that subset, and so on and so forth.

# I do have a concern here, though. Alex is investigating min and max values for life expectancy.  
# However, it's not clear where in the years recorded that these min and max values occur. It seems like  
# Alex is assuming that the min value occurs during the first year on record while the max occurs  
# during the last year on record, and that progress had been close to, if not definitely, linear.  
# However, it is possible that there could've been some wild fluctuations over the course of the years  
# that could cause the life expectancy to start at one value, dip for one reason, spike upwards for  
another,  
# and then return to the starting value. It may be worthwhile to do my own EDA with this dataset later  
on;  
# for now, it might make more sense to just follow along with what Alex is doing in the video.

# Look at the average life expectancy of all countries recorded by year.  
# Round the averages to two decimal places in order for the numbers to look  
# cleaner. Also, filter out rows where the life expectancy for any given year is zero.  
SELECT Year, ROUND(AVG(`Life expectancy`), 2)  
FROM world\_life\_expectancy  
WHERE 0 != `Life expectancy`  
GROUP BY Year  
ORDER BY Year ASC  
;

# Overall, it looks like there's a persistent increase of average life expectancy over the years. Worth  
noting  
# is that in addition to observing the central tendency, it may be worthwhile to consider spread around  
that  
# central tendency (standard deviation in this case).

# Review the life expectancy table to see if there's anything else interesting to look at.

```
SELECT *  
FROM world_life_expectancy  
;
```

# Alex says that correlations between columns can be calculated in SQL,

# and he would like to look into any correlations between life expectancy and GDP.

```
SELECT Country, `Life expectancy`, GDP  
FROM world_life_expectancy  
;
```

# Look at the average life expectancy for each country.

# Round these averages to one decimal place this time because the original numbers

# are rounded to one decimal place. Look at the average GDP for each country, and

# round these averages to one decimal place. I'm using aliases for these averages

# that are different from what Alex uses in order to better describe what these values represent.

# Also, this time around Alex doesn't filter out rows where the life expectancy

# for any given year is zero. I'll do the same as well for the sake of trying to match his output; however,

# I've kept a WHERE statement as a comment to reflect that I think that it might be worthwhile

# to filter out rows where life expectancy is zero.

```
SELECT Country, ROUND(AVG(`Life expectancy`), 1) AS Avg_Life_Expectancy, ROUND(AVG(GDP), 1) AS  
Avg_GDP  
FROM world_life_expectancy  
# WHERE 0 != `Life expectancy`  
GROUP BY Country  
;
```

# In the output of the above query, there are some countries with an average GDP of zero.

# Moreover, the Cook Islands have zeros for both average life expectancy and average GDP.

# Repeat the above query, but now order the average life expectancy values in ascending order.

```
SELECT Country, ROUND(AVG(`Life expectancy`), 1) AS Avg_Life_Expectancy, ROUND(AVG(GDP), 1) AS  
Avg_GDP  
FROM world_life_expectancy  
# WHERE 0 != `Life expectancy`  
GROUP BY Country  
ORDER BY Avg_Life_Expectancy ASC  
;
```

# Alex notes that the countries with zeros for average life expectancy are small countries

# and may not be reporting their life expectancy figures, so their fields are filled with zeros instead.

# However, the presence of zeros could be a problem when performing EDA

# Repeat the above query, but this time filter retain only the rows that have average life expectancy  
values

# greater than zero and an average GDP greater than zero.

```

SELECT Country, ROUND(AVG(`Life expectancy`), 1) AS Avg_Life_Expectancy, ROUND(AVG(GDP), 1) AS
Avg_GDP
FROM world_life_expectancy
# WHERE 0 != `Life expectancy`
GROUP BY Country
HAVING 0 < Avg_Life_Expectancy
      AND 0 < Avg_GDP
ORDER BY Avg_Life_Expectancy ASC
;

```

# Repeat the above query but this time order by average GDP in ascending order.

```

SELECT Country, ROUND(AVG(`Life expectancy`), 1) AS Avg_Life_Expectancy, ROUND(AVG(GDP), 1) AS
Avg_GDP
FROM world_life_expectancy
# WHERE 0 != `Life expectancy`
GROUP BY Country
HAVING 0 < Avg_Life_Expectancy
      AND 0 < Avg_GDP
ORDER BY Avg_GDP ASC
;

```

# Review the average life expectancy of all countries recorded by year.

# Round the averages to two decimal places in order for the numbers to look

# cleaner. Also, filter out rows where the life expectancy for any given year is zeros.

```

SELECT Year, ROUND(AVG(`Life expectancy`), 2)
FROM world_life_expectancy
WHERE 0 != `Life expectancy`
GROUP BY Year
ORDER BY Year ASC
;

```

# Alex says it's helpful to save past queries in order to reuse them later on;

# however, I'm inclined to copy-and-paste older queries in order to retain

# proper accounting of my thought processes

# According to Alex, it looks like the average of the average life expectancies found in the above query  
# is about sixty-eight years.

# Repeat the query two above, to see how countries compare to this average of averages

```

SELECT Country, ROUND(AVG(`Life expectancy`), 1) AS Avg_Life_Expectancy, ROUND(AVG(GDP), 1) AS
Avg_GDP
FROM world_life_expectancy
# WHERE 0 != `Life expectancy`
GROUP BY Country
HAVING 0 < Avg_Life_Expectancy
      AND 0 < Avg_GDP
ORDER BY Avg_GDP ASC
;

```

# Alex notes that countries with lower average GDPs appear to have lower average life expectancies,  
# possibly because of suboptimal healthcare or social infrastructure in those countries.

# I've been under the impression that Alex would use some sort of built-in function to calculate correlation,  
# and to be real I feel more comfortable taking that approach. Eyeballing correlations feels risky to me  
# because human error can negatively impact analysis. I did a quick Google search, and one source states that there is  
# a function that can calculate correlation values between two columns - CORR(Y, X) - but it's not supported by MySQL  
# (source = <https://datacomy.com/sql/functions/aggregation/correlation/>). However, after further investigation  
# another source points out that MySQL can be "extended" with a package of statistical functions, including  
# the above correlation function (source = <https://sqlstat.sourceforge.net/index.html>). For now, I think I'll just  
# do my best to follow along with what Alex is doing in the video.

# Repeat the above query but sort the average GDPs in descending order.  
SELECT Country, ROUND(AVG(`Life expectancy`), 1) AS Avg\_Life\_Expectancy, ROUND(AVG(GDP), 1) AS Avg\_GDP  
FROM world\_life\_expectancy  
# WHERE 0 != `Life expectancy`  
GROUP BY Country  
HAVING 0 < Avg\_Life\_Expectancy  
AND 0 < Avg\_GDP  
ORDER BY Avg\_GDP DESC  
;

# It looks like countries with high average GDPs also have high average life expectancies.

# Alex says that in order to confirm that there is a correlation between these two columns,  
# other tools like Tableau or Power BI can generate visualizations based off of this data. I suppose  
# that there really are limitations to what SQL (or at least MySQL) can do.

# Attempt to place average life expectancies and average GDP values into bins, and then  
# filter countries based off of those bins. First though look at the full life expectancy table  
# in order to guesstimate a halfway point of the GDP values for binning purposes.  
SELECT \*  
FROM world\_life\_expectancy  
ORDER BY GDP  
;

# Assign numerical bin tags based off of that guesstimated halfway point and then add up the number of items per bin  
# to determine how many high GDP rows there are  
SELECT

```
SUM(CASE
    WHEN 1500 <= GDP THEN 1
    ELSE 0
END) AS High_GDP_Count
FROM world_life_expectancy
ORDER BY GDP
;
```

# Run another query that's similar to the one above but with an extra column  
 # similar to the High\_GDP\_Count column.  
 # The NULL value in the CASE statement inside of the AVG() function is important  
 # for not bringing down the average of the life expectancies for the high GDP countries.  
 # I used a column name for the average that's different from the column name used by Alex  
 # in order for me to better understand what's going on.  
 # Reminder: The average value in the output of the query below is the average for the  
 # rows labeled as high GDP, not for all of the countries.

```
SELECT
SUM(CASE WHEN 1500 <= GDP THEN 1 ELSE 0 END) AS High_GDP_Count,
AVG(CASE WHEN 1500 <= GDP THEN `Life expectancy` ELSE NULL END) AS
Avg_High_GDP_Life_expectancy
FROM world_life_expectancy
ORDER BY GDP
;
```

# Repeat the above query but now also take into consideration the low GDP countries.

```
SELECT
SUM(CASE WHEN 1500 <= GDP THEN 1 ELSE 0 END) AS High_GDP_Count,
AVG(CASE WHEN 1500 <= GDP THEN `Life expectancy` ELSE NULL END) AS
Avg_High_GDP_Life_expectancy,
SUM(CASE WHEN GDP <= 1500 THEN 1 ELSE 0 END) AS High_GDP_Count,
AVG(CASE WHEN GDP <= 1500 THEN `Life expectancy` ELSE NULL END) AS
Avg_Low_GDP_Life_expectancy
FROM world_life_expectancy
ORDER BY GDP
;
```

# Alex is using his understanding of the data derived from previous queries to look around  
 # and find anything interesting. He also says that the work done here with life expectancy and GDP  
 # can be applied to just about any column.

# I do have a concern about the output of the above query. It is my understanding that some countries  
 # have 0 as their GDP value, so they'd be considered low GDP countries by one of the CASE statements.

# Consider the following query

```
SELECT Country, `Status`, AVG(`Life expectancy`)
FROM world_life_expectancy
WHERE GDP = 0
GROUP BY Country, `Status`
ORDER BY `Status` ASC, Country ASC
```

;

# Based on the output of the above query, the collection of countries with GDP values of zero includes both developed countries

# (e.g., the UK and the US) and developing countries (e.g., Iraq and Iran). Interestingly, the Republic of Korea

# (i.e., South Korea) is classified as developing.

```
SELECT Country, `Year`, `Status`  
FROM world_life_expectancy  
WHERE Country = 'Republic of Korea'
```

;

# Based off of the above query, the Republic of Korea had been a developing country for several years even though now

# I'd consider this country to be an economic and cultural powerhouse.

# Alas, my own curiosity and growing understanding of this data is taking me along my own EDA,

# so I might as well for now continue following along with Alex.

# Alex would like to investigate average life expectancy for the different country statuses. Currently there

# are only two statuses, developing and developed.

```
SELECT `Status`, ROUND(AVG(`Life expectancy`), 1)  
FROM world_life_expectancy  
GROUP BY `Status`
```

;

# Alex would like to investigate the number of countries classified as either

# developing or developed

```
SELECT `Status`, COUNT(DISTINCT Country)  
FROM world_life_expectancy  
GROUP BY `Status`
```

;

# Based off of the result of the above query, Alex figures that the average life expectancy

# can be skewed in favor of the developed countries because they have so few countries to bring their value down.

# I have an understanding of what he's trying to argue, but at the same time I also have a gut feeling

# that there might be a flaw in the logical progression of his argument, either explicitly mentioned or implied.

# Maybe I'll be able to better articulate the rationale of my gut instinct later on, but perhaps the development status

# has more of an impact on average life expectancy rather than number of members in each status, and perhaps

# the number members in status has more of an impact on the distribution around that mean (i.e., the standard deviation)

# than on the average itself.

# Alex would like to combine the two queries above into one.

```
SELECT `Status`, COUNT(DISTINCT Country), ROUND(AVG(`Life expectancy`), 1)
FROM world_life_expectancy
GROUP BY `Status`
;
```

# Alex would now like to look into the BMI.

# First look at the table once more to determine what to compare the BMI to.

```
SELECT *
FROM world_life_expectancy
;
```

# Alex has decided to look at average life expectancy and average BMI over the recorded years

# and group those values by country. I've opted to use column headers that better help me

# keep track of what's happening.

# Alex points out that while it can be worthwhile to rewrite code from scratch for learning purposes,

# in the real world he's more inclined to just reuse code that's been written already in order

# to save some time.

```
SELECT Country, ROUND(AVG(`Life expectancy`), 1) AS Avg_Life_Expectancy, ROUND(AVG(BMI), 1) AS
Avg_BMI
FROM world_life_expectancy
# WHERE 0 != `Life expectancy`
GROUP BY Country
HAVING 0 < Avg_Life_Expectancy
      AND 0 < Avg_BMI
ORDER BY Avg_BMI DESC
;
```

# Alex is initially shocked by these BMI numbers because 60+ just seems so high.

# At the same time, I'm thinking that there are no units displayed in the column names,

# and as far as I can recall Alex hasn't mentioned a data dictionary, and if such a

# dictionary exists, then maybe some of the BMI values can make sense.

# Next, Alex tries to find any correlation between average life expectancy and

# average, and based off of the output from the above query it looks like countries

# with above-average life expectancy values have high BMI values. Alex admits that he's not a doctor,

# but as far as he understands, high BMI values relates to higher likelihoods of heart attacks

# occurring. Again, it's not clear what the units of BMI are here, especially since not every country

# will use the same units of measurement. A quick Google search on units used for BMI shows that

# according to Wikipedia, BMI can be expressed in kg/meter-squared, and that when converting to pounds and feet,

# a conversion factor of 4.88 needs to be used (source =

[https://en.wikipedia.org/wiki/Body\\_mass\\_index](https://en.wikipedia.org/wiki/Body_mass_index)).

# So yeah...a data dictionary with details about the data here can be really helpful, ya know?

# Alex reasons something to the effect that countries that are developed and have high GDP values are likely

# to have well-fed populations which can contribute to high BMI values.



```

# Run the same query as above but this time sort the average BMI values in ascending order
# to look at countries with low BMI values.
SELECT Country, ROUND(AVG(`Life expectancy`), 1) AS Avg_Life_Expectancy, ROUND(AVG(BMI), 1) AS
Avg_BMI
FROM world_life_expectancy
# WHERE 0 != `Life expectancy`
GROUP BY Country
HAVING 0 < Avg_Life_Expectancy
      AND 0 < Avg_BMI
ORDER BY Avg_BMI ASC
;

```

# So with the exception of Viet Nam, some countries with low average BMI values  
 # also have lower than average life expectancy values. This could be a result of  
 # limited access to food or maybe an inclination to just eat less.

# At any rate, any findings here can be expanded upon with visualization tools  
 # in future data analyses.

```

# Look at the the whole life expectancy table again
SELECT *
FROM world_life_expectancy
;

```

# Alex would like to look at adult mortality to see how many people are dying each year  
 # in a country and how that count relates to life expectancy.

# I'm already thinking about the criteria for adulthood because  
 # the age of majority can vary from country to country, and some teens might  
 # be excluded from the calculation for adult mortality on a country-by-country basis.  
 # Also, even if an age is established as a universal threshold for adulthood for all countries,  
 # neither infant deaths nor child deaths will necessarily be included in adult mortality  
 # calculations. Moreover, it is my understanding that average life expectancy can be brought down  
 # by infant or child deaths, so the exclusion of some data in one calculation and inclusion  
 # of that same data in another calculation can make comparing those calculations  
 # questionable.

# At any rate, Alex mentions a rolling total for his next analysis,  
 # which will showcase a sum of the current value and the sum of all the  
 # values above that current value. It will involve a window function,  
 # which will include a partition over a subset of the overall table.  
 # Here is a source of information that I found with a Google search on rolling sums  
 # in SQL: <https://medium.com/@mattdamberg/calculating-running-total-in-sql-850c5b072513>  
 SELECT Country,  
 Year,  
 `Life expectancy`,  
 `Adult Mortality`,

```
SUM(`Adult Mortality`) OVER(PARTITION BY Country ORDER BY `Year`) AS Rolling_Total
FROM world_life_expectancy
;
```

# Alex suspects that the fact that only three adults died in Afghanistan in 2009  
# is a data quality issue.

# Repeat the above query once more, but try to focus on the data associated with  
# the United States of America.  
# Note: With the query that Alex will end up using,  
# some other countries with the word 'United' in their name can show up.

```
SELECT Country,
Year,
`Life expectancy`,
`Adult Mortality`,
SUM(`Adult Mortality`) OVER(PARTITION BY Country ORDER BY `Year`) AS Rolling_Total
FROM world_life_expectancy
WHERE Country LIKE '%United%'
;
```

# Out of curiosity, I looked up the definition of adult mortality,  
# and it looks like it is a count per 1,000 persons.  
# (Source = <https://datahelpdesk.worldbank.org/knowledgebase/articles/114956-what-is-the-definition-of-adult-mortality>)  
# With that said, I'm not sure if applying the rolling total over adult mortality  
# really makes sense. Nevertheless, I'll follow along with Alex over the course  
# of this video.

# Review the whole life expectancy table to see if there's a column for total population for each country.

```
SELECT *
FROM world_life_expectancy
;
```

# No such column exists in the table, but Alex says that that data might be something that would have to  
be  
# acquired from online and then joined to the world life expectancy table. If country population could be  
# factored in, Alex says that he'd like to see country population, life expectancy, and adult mortality  
together.

# By the end of the video, a lot of different things had been observed. Alex recounts all of the work that  
had been done,  
# and he points out that even more work can be done with what's present in the table. I'd imagine even  
more  
# work can be done if extra data is merged with the life expectancy table.

# Worth noting here is that somewhere along the way, I've developed a habit of putting a grave accent (`)  
around

# column name if they happen to be the same as SQL keywords. That way, I make it explicitly clear  
# that I want to work with a column name and not a SQL keyword, ya know?